Intro to Data Science

What R we doing?

Prof. Bisbee

Due Date: 2023-09-04

Welcome to Data Science!

Today, we'll be working on getting you set up with the tools you will need for this class. Once you are set up, we'll do what we're here to do: analyze data!

Here's what we need to get done today:

1. Getting started with R
2. Getting started with RStudio
3. Analyze Data
4. <....>
5. Profit!

Introductions

We need two basic sets of tools for this class. We will need `R` to analyze data. We will need `RStudio` to help us interface with R and to produce documentation of our results.

Installing R

R is going to be the only programming language we will use. R is an extensible statistical programming environment that can handle all of the main tasks that we'll need to cover this semester: getting data, analyzing data and communicating data analysis.

If you haven't already, you need to download R here: https://cran.r-project.org/ (https://cran.r-project.org/).

Installing RStudio

When we work with R, we communicate via the command line. To help automate this process, we can write scripts, which contain all of the commands to be executed. These scripts generate various kinds of output, like numbers on the screen, graphics or reports in common formats (pdf, word). Most programming languages have several **I** ntegrated **D** evelopment **E** nvironments (IDEs) that encompass all of these elements (scripts, command line interface, output). The primary IDE for R is RStudio.

If you haven't already, you need to download RStudio here: https://rstudio.com/products/rstudio/download/ (https://rstudio.com/products/rstudio/download/). You need the free RStudio desktop version.

Accessing Files and Using Directories

In each class, we're going to include some code and text in one file, and data in another file. You'll need to download both of these files to your computer. You need to have a particular place to put these files. Computers are organized using named directories (sometimes called folders). Don't just put the files in your Downloads directory. One common

solution is to created a directory on your computer named after the class: `ds_1000`. Each time you access the files, you'll want to place them in that directory.

Yes We Code! Running R Code

We're going to grab some data that's part of the college scorecard (https://collegescorecard.ed.gov/data/documentation/) and do a bit of analysis on it.

.Rmd Set Up

Open `RStudio`, then create a new `.Rmd` file. To do this, click on `File → New File → R Markdown....`
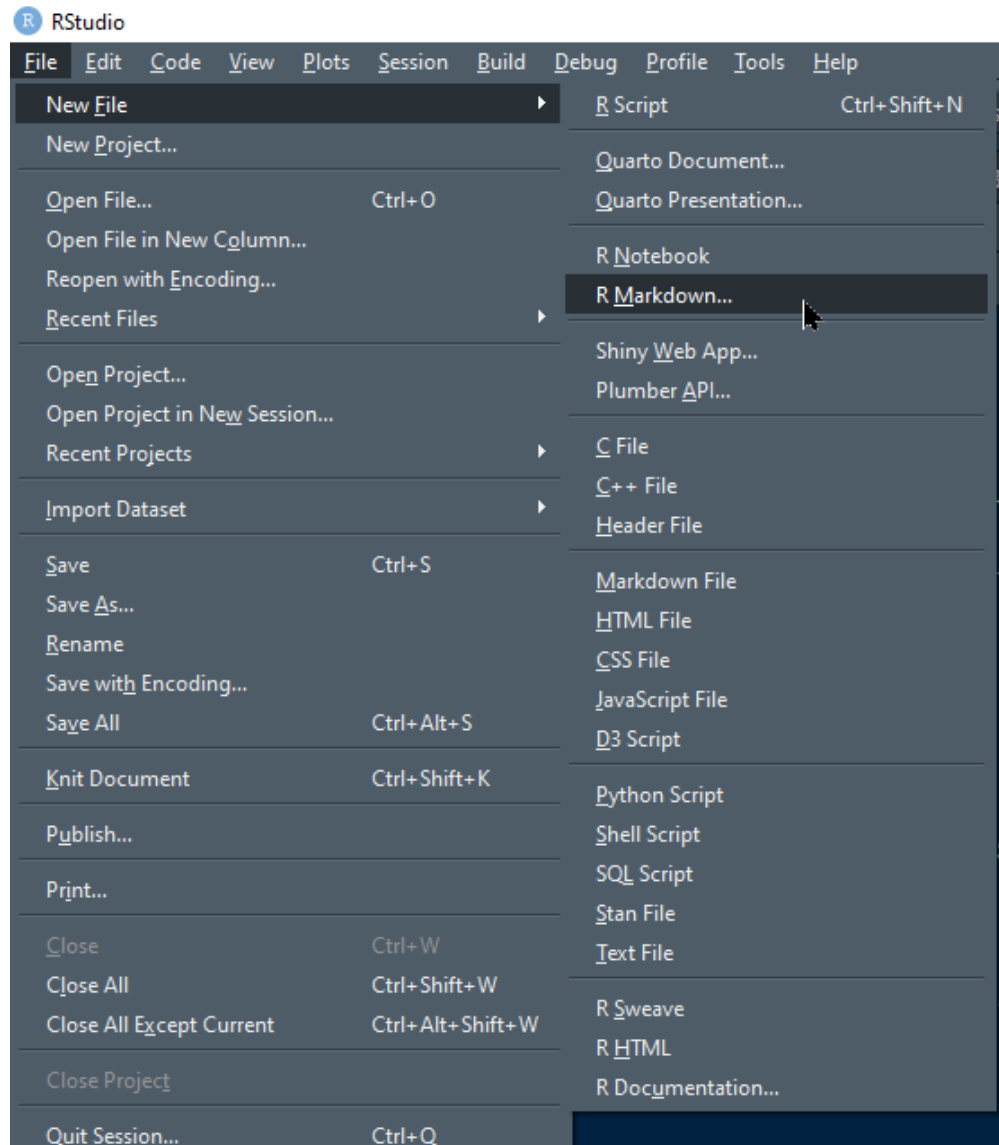
You will then be asked to determine a bunch of settings for this `.Rmd` document. For example, you can choose whether you want to create a "Document", "Presentation", "Shiny", or "From Template" on the left. You can set the "Title:" "Author:" and "Date:" on the top-right. And you can choose the "Default Output Format:" to be either "HTML", "PDF", or "Word". You should **not change any of these settings**. Their defaults ("Document", "Untitled", "[Your name]", "[Today's Date]", and "HTML") are sufficient. Just click "OK".
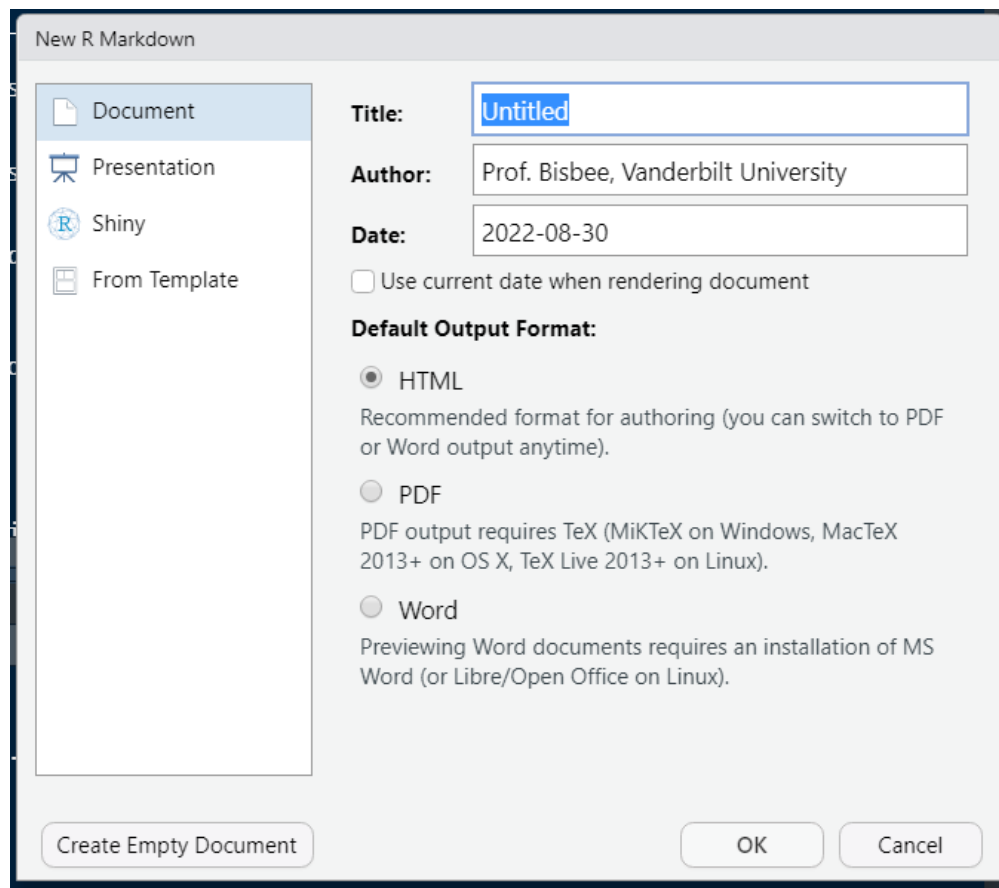
Copy the raw code from the `Intro_Data_Science_hw.Rmd` (https://github.com/jbisbee1/DS1000_F2023/blob/main/Lectures/1_Intro_Data_Science/code/Intro_Data_Science_hw.Rmd) file by clicking on the copy button as shown in the image below.
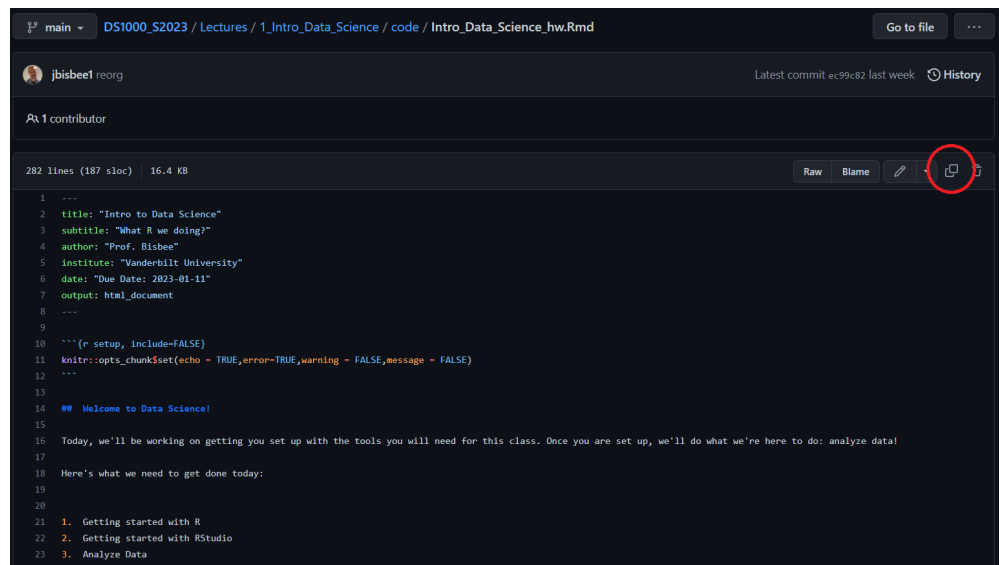
Finally, replace the default code in your R Markdown file with the copied code from the GitHub!

If viewing this as an html file, you can view this gif for more help!
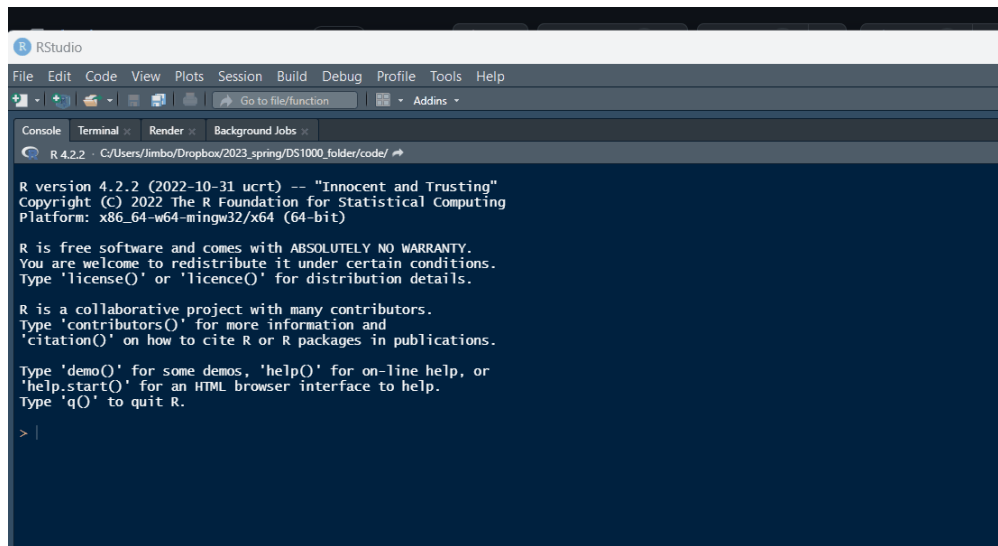
.Rmd Files

.Rmd files will be the only file format we work in this class. .Rmd files contain three basic elements:

1. Script that can be interpreted by R.
2. Output generated by R, including tables and figures.
3. Text that can be read by humans.

From a .Rmd file you can generate html documents, pdf documents, word documents, slides . . . lots of stuff. All class notes will be in .Rmd. Most assignments will be turned in as .Rmd files, and the guided exercise we'll have you do? You guessed it, .Rmd.

In the `.Rmd` file you'll notice that there are three open single quotes in a row, like so: ``` ``` ``` This indicates the start of a "code chunk" in our file. The first code chunk that we load will include a set of programs that we will need all semester long.

Outputting results

I like to see results in the Console. By default Rstudio will output results from an Rmd file inline– meaning in the document itself. To change this, go to Tools–>global Options–>R Markdown, and uncheck the box for "show output inline for all Rmarkdown documents."

Using R Libraries

When we say that R is extensible, we mean that people in the community can write programs that everyone else can use. These are called "packages." In these first few lines of code, I load a set of packages using the library command in R. The set of packages, called `tidyverse` were written by Hadley Wickham and others and play a key role in his book. To install this set of packages, simply type in `install.packages("tidyverse")` at the R command prompt. Alternatively, you can use the "Packages" pane in the lower right hand corner of your Rstudio screen. Click on Packages, then click on install, then type in "tidyverse."

To run the code below in R, you can:

- Press the "play" button next to the code chunk
- In OS X, place the cursor in the code chunk and hit `CMD+RETURN`
- In Windows, place the cursor in the code chunk and hit `CTRL+RETURN`

```
## Get necessary libraries-- won't work the first time, because you need to install them!
# install.packages("tidyverse")  # Uncomment this to install
library(tidyverse)
```

Here's the thing about packages. There's a difference between *installing* a package and *calling* a package. *Installing* means that the package is on your computer and available to use. *Calling* a package means that the commands in the package will be used in this session. A "session" is basically when R has been opened up on your computer. As long as R/Rstudio are open and running, the session is active.

It's a good practice to shutdown R/Rstudio once you're no longer working on it, and then to restart it when you begin working again. Otherwise, the working environment can get pretty crowded with data and packages.

Loading Datasets

Now we're ready to load in data. The data frame will be our basic way of interacting with everything in this class. The `sc_debt.Rds` (found here:
https://github.com/jbisbee1/DS1000_F2023/blob/main/Lectures/1_Intro_Data_Science/data/sc_debt.Rds
(https://github.com/jbisbee1/DS1000_F2023/blob/main/Lectures/1_Intro_Data_Science/data/sc_debt.Rds)) data frame contains information from the college scorecard on different colleges and universities.

However, we first need to make sure that R is looking in the right place. Because we keep our files in a particular directory, we need to point R to the right place. If this `.Rmd` file is stored in a folder called `code`, and there is another folder called `data` in the same directory as `code`, we would need to tell `R`:

1. Leave the `code` and go back to the parent directory: `../`
2. Go into the `data` folder: `data/`
3. Open the dataset `sc_debt.Rds`: `read_rds("../data/sc_debt.Rds")`

```
df<-read_rds("../data/sc_debt.Rds")

# NB: you can also open directly from GitHub by copying the link to the raw data
df <- read_rds('https://github.com/jbisbee1/DS1000_F2023/raw/main/Lectures/1_Intro_Data_Scien
ce/data/sc_debt.Rds')
```

You'll notice that the code above starts with `df`. This is just an arbitrary name for an object. You could name it `dat` or `raw` or `debt` or whatever you want. Then there's an arrow `<-`. This is an assignment operator. Then there's a function, `readRDS`, with parentheses, and an argument "sc_debt.Rds". Here's how to think about this.

- Functions in R always have arguments within parentheses. This function. `readRDS` opens a type of data– rds data. This function has one argument which is the name of the file I want to open.
- Assignment operators take the result of a function and assign it to an object name.
- Objects in R store information locall so that it can be accessed again.

So the command above says "use `readRDS` to open the file"sc_debt.Rds" and assign the result to the object `df`.

Let's take a quick look at the object `df`

```
df
```

```
## # A tibble: 2,546 × 16
##    unitid instnm    stabbr grad_…¹ control region preddeg opena…² adm_r…³ ccbasic
##     <int> <chr>     <chr>    <int> <chr>   <chr>  <chr>     <int>   <dbl>   <int>
##  1 100654 Alabama…  AL       33375 Public  South… Bachel…       2   0.918      18
##  2 100663 Univers…  AL       22500 Public  South… Bachel…       2   0.737      15
##  3 100690 Amridge…  AL       27334 Private South… Associ…       1  NA          20
##  4 100706 Univers…  AL       21607 Public  South… Bachel…       2   0.826      16
##  5 100724 Alabama…  AL       32000 Public  South… Bachel…       2   0.969      19
##  6 100751 The Uni…  AL       23250 Public  South… Bachel…       2   0.827      15
##  7 100760 Central…  AL       12500 Public  South… Associ…       1  NA           2
##  8 100812 Athens …  AL       19500 Public  South… Bachel…      NA  NA          22
##  9 100830 Auburn …  AL       24826 Public  South… Bachel…       2   0.904      18
## 10 100858 Auburn …  AL       21281 Public  South… Bachel…       2   0.807      15
## # … with 2,536 more rows, 6 more variables: sat_avg <int>,
## #   md_earn_wne_p6 <int>, ugds <int>, costt4_a <int>, selective <dbl>,
## #   research_u <dbl>, and abbreviated variable names ¹grad_debt_mdn, ²openadmp,
## #   ³adm_rate
```

This is just the first part of the data frame. All data frames have the exact same structure. Each row is a case. In this example, each row is a college. Each column is a characteristics of the case, what we call a variable. Let's use the `names` command to see what variables are in the dataset.

```
names(df)
```

```
##  [1] "unitid"        "instnm"        "stabbr"        "grad_debt_mdn"
##  [5] "control"       "region"        "preddeg"       "openadmp"
##  [9] "adm_rate"      "ccbasic"       "sat_avg"       "md_earn_wne_p6"
## [13] "ugds"          "costt4_a"      "selective"     "research_u"
```

It's hard to know what these mean without some more information. We usually use a codebook to get more information about a dataset. Because we use very short names for variables, it's useful to have some more information (fancy name: metadata) that tells us about those variables. Below you'll see the `R` name for each variable next to a description of each variable.

Name	Definition
unitid	Unit ID
instnm	Institution Name
stabbr	State Abbreviation
grad_debt_mdn	Median Debt of Graduates
control	Control Public or Private
region	Census Region
preddeg	Predominant Degree Offered: Associates or Bachelors
openadmp	Open Admissions Policy: 1= Yes, 2=No,3=No 1st time students
adm_rate	Admissions Rate: proportion of applications accepted

Name	Definition
ccbasic	Type of institution– see here (https://data.ed.gov/dataset/9dc70e6b-8426-4d71-b9d5-70ce6094a3f4/resource/658b5b83-ac9f-4e41-913e-9ba9411d7967/download/collegescorecarddatadictionary_01192021.xlsx)
selective	Institution admits fewer than 10 % of applicants, 1=Yes, 0=No
research_u	Institution is a research university 1=Yes, 0=No
sat_avg	Average Sat Scores
md_earn_wne_p6	Average Earnings of Recent Graduates
ugds	Number of undergraduates
costt4a	Average cost of attendance (tuition-grants)

Looking at datasets

We can also look at the whole dataset using View. Just delete the `#` sign below to make the code work. That `#` sign is a comment in R code, which indicates to the computer that everything on that line should be ignored. To get it to run, we need to drop the `#`.

```
View(df)
```

You'll notice that this data is arranged in a rectangular format, with each row showing a different college, and each column representing a different characteristic of that college. Datasets are always structured this way— cases (or units) will form the rows, and the characteristics of those cases– or variables— will form the columns. Unlike working with spreadsheets, this structure is always assumed for datasets.

Filter, Select, Arrange

In exploring data, many times we want to look at smaller parts of the dataset. There are three commands we'll use today that help with this.

- `filter` selects only those cases or rows that meet some logical criteria.

- `select` selects only those variables or columns that meet some criteria

- `arrange` arranges the rows of a dataset in the way we want.

For more on these, please see this vignette (https://cran.rstudio.com/web/packages/dplyr/vignettes/introduction.html).

Let's grab just the data for Vanderbilt, then look only at the average test scores and admit rate. We can use filter to look at all of the variables for Vanderbilt:

```
df%>%
   filter(instnm=="Vanderbilt University")
```

```
## # A tibble: 1 × 16
##   unitid instnm    stabbr grad_…¹ control region preddeg opena…² adm_r…³ ccbasic
##    <int> <chr>     <chr>    <int> <chr>   <chr>  <chr>     <int>   <dbl>   <int>
## 1 221999 Vanderbi… TN       14962 Private South… Bachel…       2  0.0912      15
## # … with 6 more variables: sat_avg <int>, md_earn_wne_p6 <int>, ugds <int>,
## #   costt4_a <int>, selective <dbl>, research_u <dbl>, and abbreviated variable
## #   names ¹grad_debt_mdn, ²openadmp, ³adm_rate
```

What's that weird looking `%>%` thing? That's called a pipe. This is how we chain commands together in R. Think of it as saying "and then" to R. In the above case, we said, take the data *and then* filter it to be just the data where the institution name is Vanderbilt University.

The command above says the following:

Take the dataframe `df` *and then* filter it to just those cases where `instnm` is equal to "Vanderbilt University." Notice the "double equals" sign, that's a logical operator asking if `instnm` is equal to "Vanderbilt University."

Many times, though we don't want to see everything, we just want to choose a few variables. `select` allows us to select only the variables we want. In this case, the institution name, its admit rate, and the average SAT scores of entering students.

```
df%>%
  filter(instnm=="Vanderbilt University")%>%
  select(instnm,adm_rate,sat_avg)
```

```
## # A tibble: 1 × 3
##   instnm                adm_rate sat_avg
##   <chr>                    <dbl>   <int>
## 1 Vanderbilt University   0.0912    1515
```

`filter` takes logical tests as its argument. The code `insntnm=="Vanderbilt University"` is a logical statement that will be true of just one case in the dataset– when institution name is Vanderbilt University. The `==` is a logical test, asking if this is equal to that. Other common logical and relational operators for R include

- `>`, `<` : greater than, less than
- `>=`, `<=` : greater than or equal to, less than or equal to
- `!` :not, as in `!=` not equal to
- `&` AND
- `|` OR

Next, we can use `filter` to look at colleges with low admissions rates, say less than 10% (or .1 in the proportion scale used in the dataset).

```
df%>%
  filter(adm_rate<.1)%>%
  select(instnm,adm_rate,sat_avg)%>%
  arrange(sat_avg,adm_rate)%>%
  print(n=20)
```

```
## # A tibble: 25 × 3
##    instnm                                       adm_rate sat_avg
##    <chr>                                           <dbl>   <int>
##  1 Colby College                                  0.0967    1456
##  2 Swarthmore College                             0.0893    1469
##  3 Pomona College                                 0.074     1480
##  4 Dartmouth College                              0.0793    1500
##  5 Stanford University                            0.0434    1503
##  6 Northwestern University                        0.0905    1506
##  7 Columbia University in the City of New York    0.0545    1511
##  8 Brown University                               0.0707    1511
##  9 University of Pennsylvania                     0.0766    1511
## 10 Vanderbilt University                          0.0912    1515
## 11 Harvard University                             0.0464    1517
## 12 Princeton University                           0.0578    1517
## 13 Yale University                                0.0608    1517
## 14 Rice University                                0.0872    1520
## 15 Duke University                                0.076     1522
## 16 University of Chicago                          0.0617    1528
## 17 Massachusetts Institute of Technology          0.067     1547
## 18 California Institute of Technology             0.0642    1557
## 19 Saint Elizabeth College of Nursing            0             NA
## 20 Yeshivat Hechal Shemuel                       0             NA
## # … with 5 more rows
```

Now let's look at colleges with low admit rates, and order them using `arrange` by SAT scores (`-sat_avg` gives descending order).

```
df%>%
  filter(adm_rate<.1)%>%
  select(instnm,adm_rate,sat_avg)%>%
  arrange(-sat_avg)
```

```
## # A tibble: 25 × 3
##    instnm                                       adm_rate sat_avg
##    <chr>                                           <dbl>   <int>
##  1 California Institute of Technology             0.0642    1557
##  2 Massachusetts Institute of Technology          0.067     1547
##  3 University of Chicago                          0.0617    1528
##  4 Duke University                                0.076     1522
##  5 Rice University                                0.0872    1520
##  6 Yale University                                0.0608    1517
##  7 Harvard University                             0.0464    1517
##  8 Princeton University                           0.0578    1517
##  9 Vanderbilt University                          0.0912    1515
## 10 Columbia University in the City of New York    0.0545    1511
## # … with 15 more rows
```

And one last operation: all colleges that admit between 20 and 30 percent of students, looking at their SAT scores, earnings of attendees six years letter, and what state they are in, then arranging by state, and then SAT score.

```
df%>%
  filter(adm_rate>.2&adm_rate<.3)%>%
  select(instnm,sat_avg,grad_debt_mdn,stabbr)%>%
  arrange(stabbr,-sat_avg)%>%
  print(n=20)
```

```
## # A tibble: 37 × 4
##    instnm                                          sat_avg grad_…¹ stabbr
##    <chr>                                             <int>   <int> <chr>
##  1 Heritage Christian University                        NA      NA AL
##  2 University of California-Santa Barbara             1370   15000 CA
##  3 California Polytechnic State University-San Luis Obis…  1342   19501 CA
##  4 University of California-Irvine                    1306   15488 CA
##  5 California Institute of the Arts                     NA   27000 CA
##  6 University of Miami                                1371   17125 FL
##  7 Georgia Institute of Technology-Main Campus       1418   23000 GA
##  8 Point University                                   986   26000 GA
##  9 Grinnell College                                  1457   17500 IA
## 10 St Luke's College                                   NA   17750 IA
## 11 Purdue University Northwest                        1074   22250 IN
## 12 Alice Lloyd College                               1040   15838 KY
## 13 Wellesley College                                 1452   11000 MA
## 14 Boston College                                    1437   17500 MA
## 15 Brandeis University                               1434   26150 MA
## 16 Babson College                                    1362   22985 MA
## 17 Laboure College                                     NA   25229 MA
## 18 Coppin State University                            903   24076 MD
## 19 University of Michigan-Ann Arbor                  1448   17500 MI
## 20 University of North Carolina at Chapel Hill       1402   15400 NC
## # … with 17 more rows, and abbreviated variable name ¹grad_debt_mdn
```

Quick Exercise Choose a different college and two different things about that college. Have R print the output.

```
# INSERT CODE HERE
```

Summarizing Data

To summarize data, we use the `summarize` command. Inside that command, we tell R two things: what to call the new variable that we're creating, and what numerical summary we would like. The code below summarizes median debt for the colleges in the dataset by calculating the average of median debt for all institutions.

```
df%>%
  summarize(mean_debt=mean(grad_debt_mdn,na.rm=TRUE))
```

```
## # A tibble: 1 × 1
##   mean_debt
##       <dbl>
## 1    19646.
```

```
df%>%
  summarize(median_debt=median(grad_debt_mdn,na.rm=TRUE))
```

```
## # A tibble: 1 × 1
##   median_debt
##         <int>
## 1       21500
```

Quick Exercise Summarize the average entering SAT scores in this dataset.

```
# INSERT CODE HERE
```

Combining Commands

We can also combine commands, so that summaries are done on only a part of the dataset. Below, we summarize median debt for selective schools, and not very selective schools.

```
df%>%
  filter(adm_rate<.1)%>%
  summarize(mean_debt=mean(grad_debt_mdn,na.rm=TRUE))
```

```
## # A tibble: 1 × 1
##   mean_debt
##       <dbl>
## 1   16178.
```

What about for not very selective schools?

```
df%>%
  filter(adm_rate>.3)%>%
  summarize(mean_debt=mean(grad_debt_mdn,na.rm=TRUE))
```

```
## # A tibble: 1 × 1
##   mean_debt
##       <dbl>
## 1   23230.
```

Quick Exercise Calculate average earnings for schools where SAT>1200

```
# INSERT CODE HERE
```

Quick Exercise Calculate the average debt for schools that admit over 50% of the students who apply.

```
# INSERT CODE HERE
```