

Classification

Part 1

Prof. Bisbee

Vanderbilt University

Lecture Date: 2023/10/30

Slides Updated: 2023-10-29

Agenda

1. Classification
2. College Admissions

Definitions

- *Classification*: predicting the **class** of given data points via **predictive modeling**
 - *Class*: AKA targets, labels, or categories
 - *Predictive Modeling*: Approximate mapping function $f : X \rightarrow Y$
 - X : predictor variables
 - Y : outcome variable
 - f : ??

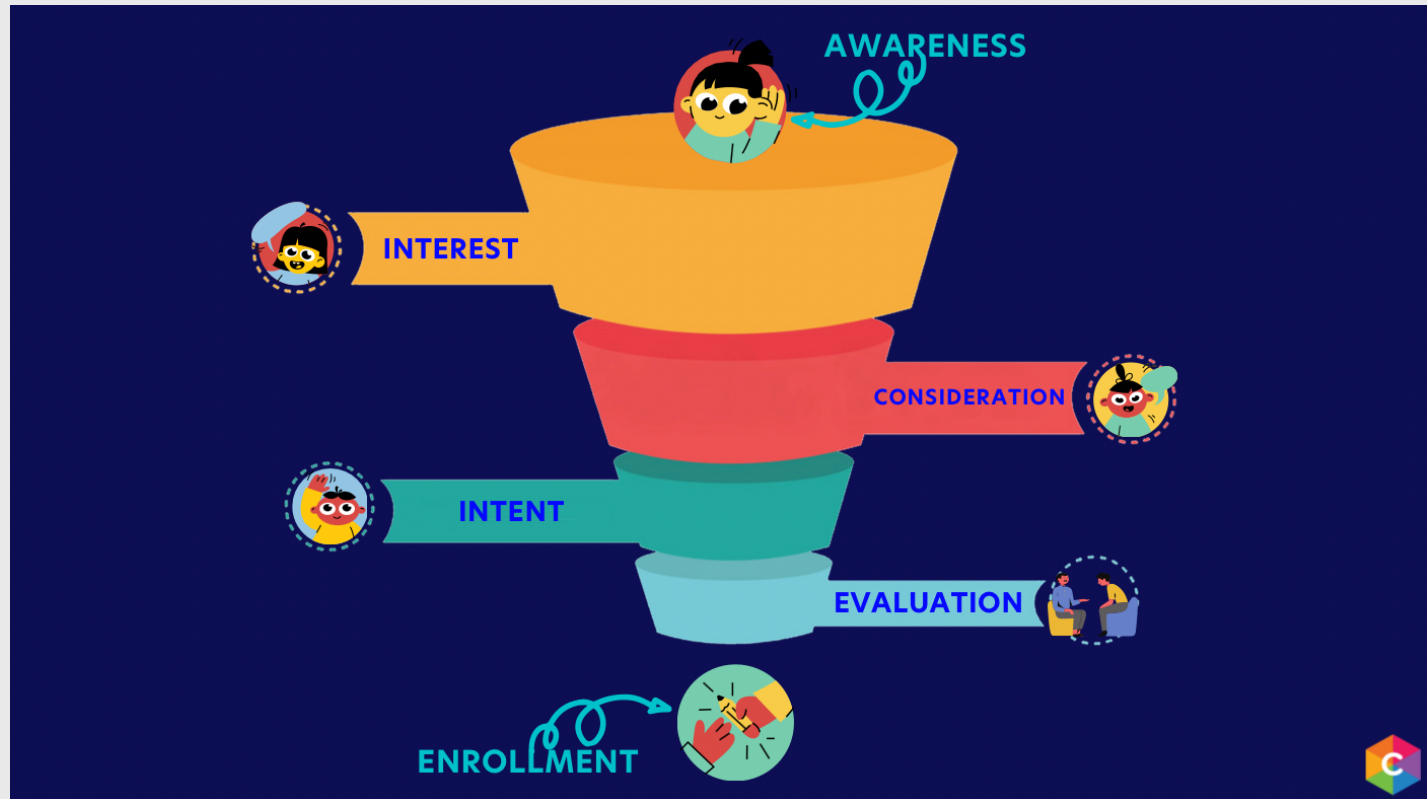
Mapping Functions

- We have already used a mapping functions!
- Linear Regression
 - $f: Y = \alpha + \beta X + \varepsilon$
- Underlying idea: X contain information about Y

It is in the Y

- If Y is continuous, we use OLS regression
- If Y is **binary**, we use "logistic" regression (AKA "logit")
 - As always, this is a **deep** area of study for those interested
- Today, using OLS for binary Y
 - Next few classes: replacing OLS regression with logit

College Admissions



- A live interactive infographic

College Admissions

- The math of college admissions

1. **Tuition** (\$)

- How they stay in business

2. **Reputation**

- Higher reputation → more **tuition**
- Based on academic qualifications

Data Science!

- This is a big industry for data scientists!
- Why?
 - If you screw this up, you lose A LOT OF MONEY
 - Too few students → not enough money to operate
 - Too many students → not enough capacity → bad reputation → not enough money
- Thus, we need people who are good at **classification**

Our Task

- Colleges hire data scientists to do more than just predict yield
- College **goals**: Increase reputation
 - Increase average SAT score to 1300
 - Admit at least 200 more students with incomes under \$50,000
- College **constraints**: Stay in operation!
 - Maintain total revenues of \$30m
 - Maintain entering class size of 1,500

How do we do this?

- Tuition discounting / targeting
 - Incentivize certain students to enroll
 - Make it cheaper for them to attend via **need-based** and **merit-based** aid

- **Need-based aid:**

$$need_{aid} = 500 + (income/1000 - 100) * -425$$

- For every \$1,000 less than \$100,000, student receives +\$425

- **Merit-based aid:**

$$merit_{aid} = 5000 + (sat/100 - 1250) * 1500$$

- For every 10 points in SAT scores above 1250, student receives extra \$1,500

So how do we do this?

- Use tuition discounting to attract certain students
 - Those with higher SAT scores
 - Those with lower incomes
- Could give aid to everyone who fits these criteria
- But this is inefficient! Giving money to those who might not attend
- Want to **target** the aid toward those most likely to attend
- Again...**prediction**

Ethics

- Is this **ethical**?
- Ethics in data science is crucial

[J]ust as the invention of the telescope revolutionized the study of the heavens, so too by **rendering the unmeasurable measurable**, the technological revolution in mobile, Web, and Internet communications has the potential to revolutionize our understanding of ourselves and how we interact.

-- Duncan Watts (2011, p. 266)

- We will return to this topic in our final meeting

The Data

```
library(tidyverse)
library(scales)
ad<-read_rds("../data/admit_data.rds")%>%ungroup()
glimpse(ad)
```

```
## Rows: 2,150
## Columns: 14
## $ ID      <chr> "0001", "0002", "0003", "0004", "0005"...
## $ income  <dbl> 289720.59, 176763.29, 81204.02, 93320....
## $ sat     <dbl> 1107.403, 1387.607, 1000.000, 1134.883...
## $ gpa     <dbl> 3.597153, 4.000000, 3.072323, 3.682776...
## $ visit   <dbl> 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0,...
## $ legacy  <dbl> 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
## $ registered <dbl> 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0,...
## $ sent_scores <dbl> 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0,...
## $ distance <dbl> 10.23279, 89.75984, 152.29961, 317.502...
## $ tuition <dbl> 45000, 45000, 45000, 45000, 45000, 450...
## $ need_aid <dbl> 0.000, 0.000, 8488.293, 3338.779, 0.00...
## $ merit_aid <dbl> 0.00, 35190.18, 0.00, 0.00, 30567.16, ...
## $ net_price <dbl> 45000.000, 9809.815, 36511.707, 41661....
## $ yield   <int> 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0,...
```

The Data

- Start with the basics:
 1. What is the unit of analysis?
 2. Which variables are we interested in?

Prediction

$$Y = \alpha + \beta_1 X_1 + \beta_2 X_2 + \cdots + \varepsilon$$

- Y : enrollment (**yield**)
- X : ??
 - In prediction, we don't care about **theory** or **research questions**
 - Just want to maximize **accuracy**...which X 's are the "best"?
- Look at univariate & conditional relationships

The Data

- Outcome Y : `yield`

```
ad %>%  
  summarise(`Yield Rate` = percent(mean(yield)))
```

```
##   Yield Rate  
## 1      68%
```

- Multivariate analysis?

Which *X*?

```
ad %>%  
  group_by(legacy) %>%  
  summarise(pr_attend = mean(yield))
```

```
## # A tibble: 2 × 2  
##   legacy pr_attend  
##   <dbl>   <dbl>  
## 1     0     0.641  
## 2     1     0.780
```

Which X ?

```
ad %>%  
  group_by(visit) %>%  
  summarise(pr_attend = mean(yield))
```

```
## # A tibble: 2 × 2  
##   visit pr_attend  
##   <dbl>   <dbl>  
## 1     0     0.644  
## 2     1     0.736
```


Which *X*?

```
ad %>%  
  group_by(sent_scores) %>%  
  summarise(pr_attend = mean(yield))
```

```
## # A tibble: 2 × 2  
##   sent_scores pr_attend  
##         <dbl>    <dbl>  
## 1           0     0.651  
## 2           1     0.805
```


Which X ?

```
ad %>%  
  ggplot(aes(x = sat, y = yield)) +  
  geom_point()
```



Which X ?

```
ad %>%  
  ggplot(aes(x = sat, y = yield)) +  
  geom_jitter()
```



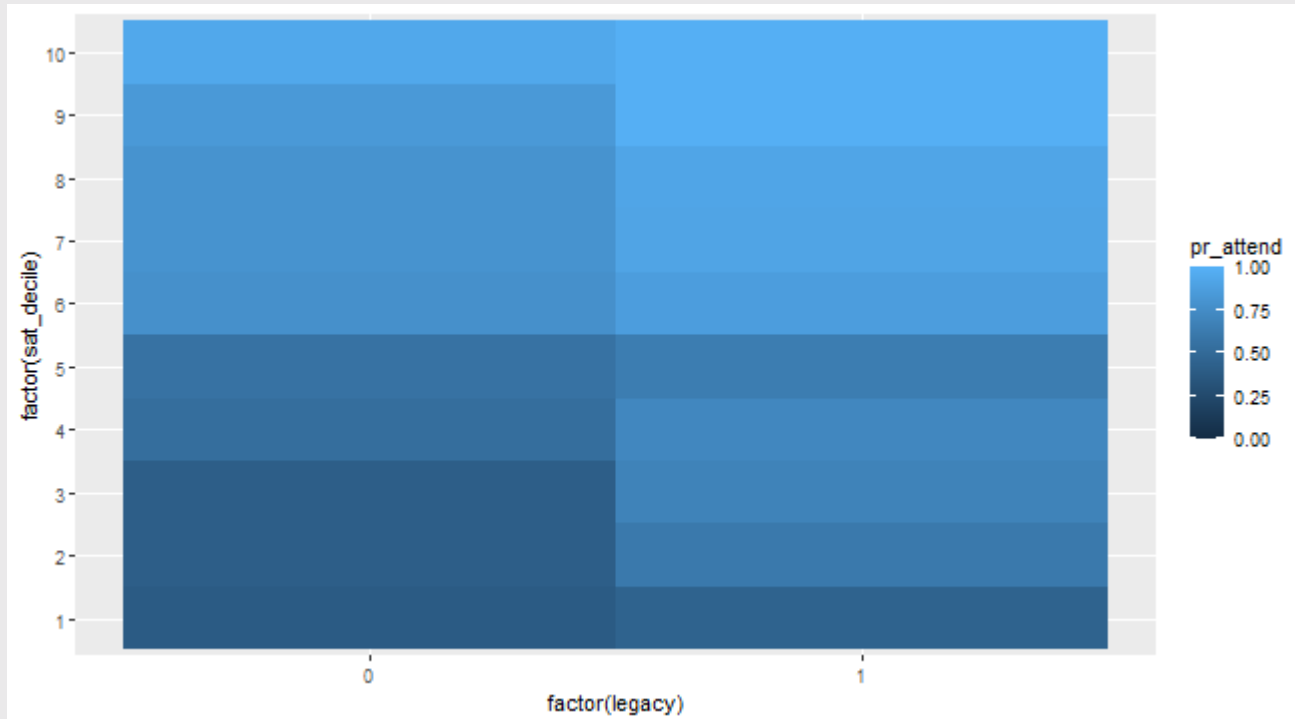
Heatmaps

- Look at 3-dimensions of data
 - Done this before by tweaking `fill`, `color`, or `size`
- `geom_tile()`: create a heatmap

```
p <- ad %>%  
  mutate(sat_decile = ntile(sat,n=10)) %>% # Bin SAT by decile (10%)  
  group_by(sat_decile,legacy) %>% # Calculate average yield by SAT &  
  Legacy  
  summarise(pr_attend = mean(yield),  
            .groups = 'drop') %>%  
  ggplot(aes(x = factor(legacy),y = factor(sat_decile), # Both x and  
            y-axes are factors  
            fill = pr_attend)) + # Fill by third dimension  
  geom_tile() + # Creates rectangles  
  scale_fill_gradient(limits = c(0,1)) # Set fill color (can do much  
  more here)
```

Heatmaps

p



Simplest Predictions

- Remember: regression is just fancier conditional means

```
ad <- ad %>%  
  mutate(sat_decile = ntile(sat,n=10)) %>% # Bin SAT by decile (10%)  
  group_by(sat_decile,legacy) %>% # Calculate average yield by SAT &  
  Legacy  
  mutate(prob_attend = mean(yield)) %>% # use mutate() instead of  
  summarise() to avoid collapsing the data  
  mutate(pred_attend = ifelse(prob_attend > .5,1,0)) %>% # If the  
  probability is greater than 50-50, predict they attend  
  ungroup()
```


Simplest Predictions

- Conditional means

```
ad %>%  
  group_by(yield, pred_attend) %>%  
  summarise(nStudents=n(), .groups = 'drop')
```

```
## # A tibble: 4 × 3  
##   yield pred_attend nStudents  
##   <int>      <dbl>      <int>  
## 1     0          0         304  
## 2     0          1         380  
## 3     1          0         210  
## 4     1          1        1256
```

Accuracy

- What is "accuracy"?
 - Proportion "correct" predictions
- For a binary outcome, "accuracy" has two dimensions
 - Proportion of correct 1s: **Sensitivity**
 - Proportion of correct 0s: **Specificity**

Accuracy


```
ad %>%
  group_by(yield) %>%
  mutate(total_attend = n()) %>%
  group_by(yield, pred_attend, total_attend) %>%
  summarise(nStudents=n(), .groups = 'drop') %>%
  mutate(prop = nStudents / total_attend)
```

```
## # A tibble: 4 × 5
##   yield pred_attend total_attend nStudents  prop
##   <int>      <dbl>      <int>      <int> <dbl>
## 1     0          0         684        304 0.444
## 2     0          1         684        380 0.556
## 3     1          0        1466        210 0.143
## 4     1          1        1466       1256 0.857
```

- Overall accuracy: $(304 + 1256) / 2150 = 73\%$

Regression

```
ad %>%  
  ggplot(aes(x = sat, y = yield)) +  
  geom_point() +  
  geom_smooth(method = 'lm')
```



Regression

- Binary outcome variable!
 - A linear regression is not the best solution
 - Predictions can exceed support of Y
- But it can still work! **linear probability model**

```
mLM <- lm(yield ~ sat + net_price + legacy, ad)
```

Linear Regression

```
require(broom) # broom package makes it easy to read regression output
```

```
## Loading required package: broom
```

```
tidy(mLM) %>% # This would be the same as summary(mLM)  
  mutate_at(vars(-term),function(x) round(x,5))
```

```
## # A tibble: 4 × 5  
##   term          estimate std.error statistic p.value  
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>  
## 1 (Intercept) -2.98      0.151    -19.7      0  
## 2 sat          0.00284   0.00012   24.2      0  
## 3 net_price    0.00001    0        14.1      0  
## 4 legacy      0.0950    0.0195    4.86      0
```

Linear Regression

```
mLM <- lm(yield ~ scale(sat) + scale(net_price) + legacy, ad)
tidy(mLM)
```

```
## # A tibble: 4 × 5
##   term                estimate std.error statistic    p.value
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)        0.654    0.0105    62.3    0
## 2 scale(sat)          0.280    0.0116    24.2  8.84e-115
## 3 scale(net_price)    0.164    0.0116    14.1  2.30e- 43
## 4 legacy              0.0950    0.0195     4.86  1.24e- 6
```

```
ad %>%
  summarise_at(vars(sat, net_price), function(x) round(sd(x), 1))
```

```
## # A tibble: 1 × 2
##   sat net_price
##   <dbl>    <dbl>
## 1  98.6    15569.
```


Evaluating Predictions

```
ad %>%  
  mutate(preds = predict(mLM)) %>%  
  mutate(predBinary = ifelse(preds > .5,1,0)) %>%  
  select(yield,predBinary,preds)
```

```
## # A tibble: 2,150 × 3  
##   yield predBinary preds  
##   <int>      <dbl> <dbl>  
## 1     1         1 0.735  
## 2     1         1 1.07  
## 3     1         0 0.245  
## 4     0         1 0.683  
## 5     1         1 0.589  
## 6     0         0 0.358  
## 7     1         1 0.559  
## 8     1         1 0.757  
## 9     0         0 0.366  
## 10    0         1 0.698  
## # i 2,140 more rows
```

Evaluating Predictions

```
ad %>%
  mutate(pred_attend = ifelse(predict(mLM) > .5,1,0)) %>%
  group_by(yield) %>%
  mutate(total_attend = n()) %>%
  group_by(yield,pred_attend,total_attend) %>%
  summarise(nStudents=n(),.groups = 'drop') %>%
  mutate(prop = nStudents / total_attend) %>%
  ungroup() %>%
  mutate(accuracy = percent(sum((yield == pred_attend)*nStudents) /
    sum(nStudents)))
```

```
## # A tibble: 4 × 6
##   yield pred_attend total_attend nStudents    prop accuracy
##   <int>      <dbl>         <int>      <int>  <dbl> <chr>
## 1     0          0           684       282 0.412  76%
## 2     0          1           684       402 0.588  76%
## 3     1          0          1466       113 0.0771 76%
## 4     1          1          1466      1353 0.923  76%
```

Evaluating Predictions

- Overall accuracy is just the number of correct predictions (either 0 or 1) out of all possible
 - Is 76% good?
 - What would the dumbest guess be? Everyone will attend! 68%
- Might also want to care about just 1s
 - **Sensitivity**: Predicted attendees / actual attendees = 92.3%
- Also might care about just 0s
 - **Specificity**: Predicted non-attendees / actual non-attendees = 41.2%

Thresholds

- Shifting the threshold for 0 or 1 prediction can matter

```
ad %>%
  mutate(pred_attend = ifelse(predict(mLM) > .4,1,0)) %>%
  group_by(yield) %>%
  mutate(total_attend = n()) %>%
  group_by(yield,pred_attend,total_attend) %>%
  summarise(nStudents=n(),.groups = 'drop') %>%
  mutate(prop = percent(nStudents / total_attend)) %>%
  ungroup() %>%
  mutate(accuracy = percent(sum((yield == pred_attend)*nStudents) /
sum(nStudents)))
```

```
## # A tibble: 4 × 6
##   yield pred_attend total_attend nStudents prop accuracy
##   <int>      <dbl>         <int>      <int> <chr> <chr>
## 1     0          0           684       176 26%   74%
## 2     0          1           684       508 74%   74%
## 3     1          0          1466        58 4%    74%
## 4     1          1          1466      1408 96%   74%
```

Thresholds

- Shifting the threshold for 0 or 1 prediction can matter

```
ad %>%
  mutate(pred_attend = ifelse(predict(mLM) > 1,1,0)) %>%
  group_by(yield) %>%
  mutate(total_attend = n()) %>%
  group_by(yield,pred_attend,total_attend) %>%
  summarise(nStudents=n(),.groups = 'drop') %>%
  mutate(prop = percent(nStudents / total_attend)) %>%
  ungroup() %>%
  mutate(accuracy = percent(sum((yield == pred_attend)*nStudents) /
sum(nStudents)))
```

```
## # A tibble: 4 × 6
##   yield pred_attend total_attend nStudents prop accuracy
##   <int>      <dbl>      <int>      <int> <chr> <chr>
## 1     0          0         684        683 99.9% 38%
## 2     0          1         684         1 0.1% 38%
## 3     1          0        1466       1342 91.5% 38%
## 4     1          1        1466        124  8.5% 38%
```


Thresholds

- Let's loop it!

```
toplot <- NULL
for(thresh in seq(0,1,by = .025)) {
  toplot <- ad %>%
    mutate(pred_attend = ifelse(predict(mLM) > thresh,1,0)) %>%
    group_by(yield) %>%
    mutate(total_attend = n()) %>%
    group_by(yield,pred_attend,total_attend) %>%
    summarise(nStudents=n(),.groups = 'drop') %>%
    mutate(prop = nStudents / total_attend) %>%
    ungroup() %>%
    mutate(accuracy = sum((yield == pred_attend)*nStudents) /
sum(nStudents)) %>%
    mutate(threshold = thresh) %>%
    bind_rows(toplot)
}
```

Thresholds

```
toplot %>%  
  mutate(metric = ifelse(yield == 1 & pred_attend == 1, 'Sensitivity',  
                          ifelse(yield == 0 & pred_attend == 0, 'Specificity', NA))) %>%  
  drop_na(metric) %>%  
  ggplot(aes(x = threshold, y = prop, color = metric)) +  
  geom_line()
```




ROC Curve

- Receiver-Operator Characteristic (ROC) Curve
- Commonly used to evaluate classification methods
 - X-axis: 1-specificity
 - Y-axis: sensitivity

```
p <- topplot %>%
  mutate(metric = ifelse(yield == 1 & pred_attend == 1, 'Sensitivity',
                        ifelse(yield == 0 & pred_attend ==
0, 'Specificity', NA))) %>%
  drop_na(metric) %>%
  select(prop, metric, threshold) %>%
  spread(metric, prop) %>%
  ggplot(aes(x = 1-Specificity, y = Sensitivity)) +
  geom_line() +
  xlim(c(0,1)) + ylim(c(0,1)) +
  geom_abline(slope = 1, intercept = 0, linetype = 'dotted') +
  ggthemes::theme_ridges()
```


ROC Curve

p



- Better models have high levels of sensitivity **and** specificity at every threshold

AUC Measure

- Area Under the Curve (AUC)
 - A single number summarizing classification performance

```
require(tidymodels)
roc_auc(data = ad %>%
  mutate(pred_attend = predict(mLM),
         truth = factor(yield, levels = c('1', '0')))) %>%
  select(truth, pred_attend), truth, pred_attend)
```

```
## # A tibble: 1 × 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 roc_auc binary      0.801
```

Party time!

- Adding more variables / trying different combinations
- **Workflow**
 1. Train models
 2. Predict models
 3. Evaluate models

Train models

```
m1 <- lm(yield ~ sat + net_price + legacy,ad)
m2 <- lm(yield ~ sat + net_price + legacy + income,ad)
m3 <- lm(yield ~ sat + net_price + legacy + income + gpa,ad)
m4 <- lm(yield ~ sat + net_price + legacy + income + gpa +
distance,ad)
m5 <- lm(yield ~ sat + net_price + legacy + income + gpa + distance +
visit,ad)
m6 <- lm(yield ~ sat + net_price + legacy + income + gpa + distance +
visit + registered + sent_scores,ad)
```

Predict models


```
toEval <- ad %>%  
  mutate(m1Preds = predict(m1),  
         m2Preds = predict(m2),  
         m3Preds = predict(m3),  
         m4Preds = predict(m4),  
         m5Preds = predict(m5),  
         m6Preds = predict(m6),  
         truth = factor(yield, levels = c('1', '0')))
```

Evaluate models

```
rocRes <- NULL
for(model in 1:6) {
  rocRes <- roc_auc(toEval,truth,paste0('m',model,'Preds')) %>%
    mutate(model = model) %>%
    bind_rows(rocRes)
}
```

Evaluate models

```
rocRes %>%  
  ggplot(aes(x = .estimate, y = reorder(model, .estimate))) +  
  geom_bar(stat = 'identity') +  
  ggthemes::theme_ridges()
```



OVERFITTING

- Cross validation to the rescue!

```
set.seed(123)
cvRes <- NULL
for(i in 1:100) {
  # Cross validation prep
  inds <- sample(1:nrow(ad),size = round(nrow(ad)*.8),replace = F)
  train <- ad %>% slice(inds)
  test <- ad %>% slice(-inds)


  # Training models
  m1 <- lm(yield ~ sat + net_price + legacy,train)
  m2 <- lm(yield ~ sat + net_price + legacy + income,train)
  m3 <- lm(yield ~ sat + net_price + legacy + income + gpa,train)
  m4 <- lm(yield ~ sat + net_price + legacy + income + gpa + distance,train)
  m5 <- lm(yield ~ sat + net_price + legacy + income + gpa + distance + visit,train)
  m6 <- lm(yield ~ sat + net_price + legacy + income + gpa + distance + visit + registered + sent_scores,train)

  # Predicting models
  toEval <- test %>%
    mutate(m1Preds = predict(m1,newdata = test),
           m2Preds = predict(m2,newdata = test),
           m3Preds = predict(m3,newdata = test),
           m4Preds = predict(m4,newdata = test),
           m5Preds = predict(m5,newdata = test),
           m6Preds = predict(m6,newdata = test),
           truth = factor(yield,levels = c('1','0')))

  # Evaluating models
  rocRes <- NULL
  for(model in 1:6) {
    rocRes <- roc_auc(toEval,truth,paste0('m',model,'Preds')) %>%
      mutate(model = model) %>%
      bind_rows(rocRes)
  }
  cvRes <- rocRes %>%
    mutate(bsInd = i) %>%
    bind_rows(cvRes)
}
```


Cross Validation AUC

```
cvRes %>%  
  ggplot(aes(x = .estimate, y = factor(reorder(model, .estimate)))) +  
  geom_boxplot() +  
  ggthemes::theme_ridges()
```



Conclusion

- Classification is just a type of prediction
 - We used linear regression
 - But there are **much** fancier algorithms out there
- Next class:
 - A *slightly* fancier algorithm: logistic regression
 - How to use the models to achieve the university's goals
- Go to Brightspace and take the **12th** quiz
 - The password to take the quiz is ####
- **Homework:**
 - Problem Set 6 (due 2023-11-03 by 11:59PM)