

Classification

Part 3

Prof. Bisbee

Vanderbilt University

Lecture Date: 2023/03/27

Slides Updated: 2023-03-26

Agenda

1. Recap of classification
2. Recap of evaluation
3. Using classifiers

Classification Recap

- Similar to **prediction**
- Difference is in the **outcome variable Y** :
 - Continuous: use regression model
 - Categorical / Binary: use classification model
- **NB:** *classification* is a type of *prediction*


Classification Recap

- Thus far, only used binary Y
 - *Can* use linear regression model
 - *Better* to use logistic regression model
- Linear regression: `lm(formula,data)`
 - Predicts "latent" measures (`predict(mLM)`)
 - I.e., willingness to attend
- Logistic: `glm(formula,data,family = binomial(link = 'logit'))`
 - Predicts probabilities (`predict(mLG,type = 'response')`)
 - I.e., probability of attending

```
require(tidyverse)
require(tidymodels)
ad <- readRDS('../data/admit_data.rds')
```


Classification Recap

```
ad %>% ggplot(aes(x = sat,y = yield)) +  
  geom_jitter(width = .01,height = .05,alpha = .25) +  
  geom_smooth(method = 'lm',color = 'black') +  
  geom_smooth(method = 'glm',color = 'red',  
              method.args = list(family = binomial(link = 'logit')))
```




Classification Recap

```
ad %>% ggplot(aes(x = income,y = yield)) +  
  geom_jitter(width = .01,height = .05,alpha = .25) +  
  geom_smooth(method = 'lm',color = 'black') +  
  geom_smooth(method = 'glm',color = 'red',  
             method.args = list(family = binomial(link = 'logit')))
```



Classification Recap

```
ad %>% ggplot(aes(x = log(income), y = yield)) +  
  geom_jitter(width = .01, height = .05, alpha = .25) +  
  geom_smooth(method = 'lm', color = 'black') +  
  geom_smooth(method = 'glm', color = 'red',  
              method.args = list(family = binomial(link = 'logit')))
```



Classification Workflow

1. Train: `mLG <- glm(formula,data,family = binomial(link = 'logit'))`
 2. Predict: `data$predY <- predict(mLG,type = 'response')`
 3. Evaluate: `roc_auc(data,truth,estimate)`
- **NB:** Evaluation stage is *DEEP*
 - Classify observations based on threshold: `ifelse(predY > 0.5,1,0)`
 - Calculate accuracy by group across thresholds
 - **Sensitivity** and **Specificity** → ROC curve → AUC
 - Should be estimated via cross validation to prevent **overfitting**
 - Make sure `truth` is ordered in **reverse**: `factor(truth,levels = c('1','0'))`

Classification Workflow

- What do we get out of all this effort?
- A single classification algorithm that we think is best

```
# Train
form <- 'yield ~ sat + legacy + visit + registered + sent_scores + income + gpa + distance + net_price'
mLG <- glm(formula = as.formula(form),
           data = ad,family = binomial(link = 'logit'))

# Predict
pred <- ad %>%
  mutate(predY = predict(mLG,type = 'response'), # NB: type = 'response' for glm()!
         truth = factor(yield,levels = c('1','0'))) # NB: reorder outcome so that 1 is first!

# Evaluate
roc_auc(data = pred,truth = 'truth',estimate = 'predY')
```

```
## # A tibble: 1 × 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 roc_auc binary      0.915
```

Classification Workflow


- But we should actually do this with **cross validation**

```
cvRes <- NULL
for(i in 1:100) {
  inds <- sample(1:nrow(ad),size = round(nrow(ad)*.8),replace = F)
  train <- ad %>% slice(inds)
  test <- ad %>% slice(-inds)

  # Train
  mLM <- lm(form,data = train)
  mLG <- glm(form,data = train,family = binomial(link = 'logit'))
  # Predict
  pred <- test %>%
  mutate(predLM = predict(mLM,newdata = test),
         predLG = predict(mLG,newdata = test,type = 'response'),
         truth = factor(yield,levels = c('1','0')))
  # Evaluate
  resLG <- roc_auc(data = pred,truth = 'truth',estimate = 'predLG') %>%
  mutate(algo = 'logit')
  resLM <- roc_auc(data = pred,truth = 'truth',estimate = 'predLM') %>%
  mutate(algo = 'linear')
  cvRes <- resLG %>% bind_rows(resLM) %>% bind_rows(cvRes)
}
```

Classification Workflow

```
cvRes %>%  
  ggplot(aes(x = .estimate, fill = algo)) +  
  geom_density(alpha = .4)
```



Using the Algorithm

- The algorithm can help us make decisions
- Recall the goals and constraints of this college:
 - **Goals:** Increase SAT to 1300
 - **Goals:** Admit 200 more students with incomes under \$50,000
 - **Constraints:** Maintain total revenues of at least \$30m
 - **Constraints:** Maintain entering class size of at least 1,466
- Tools:
 - Need-based and Merit-based aid
 - Effort for visit and registration
 - Other targeting? (ethics?)

Using the Algorithm

1. **Counterfactuals** for specific types of students

- What is the probability a given student will attend?
- `data_grid()` function from `modelr` package can help!

2. **Consulting** for changes in policy

- If we increase price, what will happen to attendance?
- Just predict on the **full data**

1. Counterfactuals: `data_grid()`

- What is the probability that this student will attend?

```
mLGFfinal <- glm(formula = as.formula(form),
                 data = ad, family = binomial(link = 'logit'))
require(modelr)
hypo_data <- ad %>%
  data_grid(legacy = 0,
            visit = 1,
            registered = 1,
            sent_scores = 1,
            sat = 1400,
            income = 95000,
            gpa = 3.9,
            distance = .1,
            net_price = 6875)

predict(mLGFfinal, newdata = hypo_data, type = 'response')
```

```
##           1
## 0.9643731
```

1. Counterfactuals: `data_grid()`

- What is the probability they will attend if we increase the price by \$10k?

```
hypo_data <- ad %>%  
  data_grid(legacy = 0,  
            visit = 1,  
            registered = 1,  
            sent_scores = 1,  
            sat = 1400,  
            income = 95000,  
            gpa = 3.9,  
            distance = .1,  
            net_price = 16875)  
  
predict(mLGFfinal, newdata = hypo_data, type = 'response')
```

```
##           1  
## 0.9392844
```

1. Counterfactuals: `data_grid()`

- Can combine with `data_grid()`

```
hypo_data <- ad %>%  
  data_grid(legacy = 0,  
            visit = 1,  
            registered = 1,  
            sent_scores = 1,  
            sat = 1400,  
            income = 95000,  
            gpa = 3.9,  
            distance = .1,  
            net_price = c(6875, 16875))  
  
predict(mLGFfinal, newdata = hypo_data, type = 'response')
```

```
##           1           2  
## 0.9643731 0.9392844
```


1. Counterfactuals: `data_grid()`

- Can use `data_grid()` to calculate "typical" values for every variable
 - **Typical:** Mean for continuous, mode for categorical

```
(hypo_data <- ad %>%  
  data_grid(.model = mLGFinal))
```

```
## # A tibble: 1 × 9  
##       sat legacy visit registered sent_scores income  gpa  
##   <dbl> <dbl> <dbl>         <dbl>         <dbl> <dbl> <dbl>  
## 1 1200.      0      0           1           0 99712.  3.79  
## # ... with 2 more variables: distance <dbl>, net_price <dbl>
```

```
predict(mLGFinal,newdata = hypo_data,type = 'response')
```

```
##           1  
## 0.8655441
```

1. Counterfactuals: `data_grid()`

- Compare otherwise typical admits who differ only in terms of GPA

```
hypo_data <- ad %>%  
  data_grid(.model = mLGFinal,  
            gpa = c(3.5, 3.9))  
  
predict(mLGFinal, newdata = hypo_data, type = 'response')
```

```
##           1           2  
## 0.7936923 0.8865843
```

1. Counterfactuals: `data_grid()`

- Multiple comparisons are possible


```
require(ggbridges)
hypo_data <- ad %>%
  data_grid(.model = mLGFinal,
            gpa = seq_range(gpa,n = 100),
            legacy = c(0,1))

toplot <- hypo_data %>%
  mutate(preds = predict(mLGFinal,newdata = hypo_data,type =
'response'))

p <- toplot %>%
  ggplot(aes(x = gpa,y = preds,color = factor(legacy))) +
  geom_line() +
  theme_ridges() +
  scale_y_continuous(labels = scales::percent) +
  labs(title = 'Predicted Attendance by GPA and Legacy',
       x = 'GPA',
       y = 'Prob(Attend)')
```

1. Counterfactuals: `data_grid()`

p



2. Consulting: full data

- Currently, admit 1466 at ~\$30.7m revenues

```
ad %>%  
  filter(yield == 1) %>%  
  summarise(totRev = dollar(sum(net_price)),  
            totAtt = n())
```

```
##           totRev totAtt  
## 1 $30,674,149    1466
```

2. Consulting: full data

- What if we increased the price for those who submit scores?

```
# Currently, 355 students who sent scores are predicted to attend  
ad %>%  
  mutate(preds = predict(mLGFfinal,type = 'response')) %>%  
  mutate(pred_class = ifelse(preds > .5,1,0)) %>%  
  count(sent_scores,pred_class)
```

```
##   sent_scores pred_class    n  
## 1           0          0   627  
## 2           0          1  1093  
## 3           1          0    75  
## 4           1          1   355
```

2. Consulting: full data

- What if we increased the price for those who submit scores?

```
# If price increases by $5k, number of predicted drops to 338
hypo <- ad %>%
  mutate(net_price = ifelse(sent_scores == 1,
                             net_price + 5000, net_price))

hypo %>%
  mutate(preds = predict(mLGFfinal, newdata = hypo, type = 'response'))
%>%
  mutate(pred_class = ifelse(preds > .5, 1, 0)) %>%
  count(sent_scores, pred_class)
```

```
##   sent_scores pred_class    n
## 1           0          0   627
## 2           0          1  1093
## 3           1          0    92
## 4           1          1   338
```

2. Consulting: full data

- What if we increased the price for those who submit scores?

```
# BUT we make more total revenue
hypo %>%
  mutate(preds = predict(mLGFfinal,newdata = hypo,type = 'response'))
%>%
  mutate(pred_class = ifelse(preds > .5,1,0)) %>%
  filter(pred_class == 1) %>%
  summarise(tot_rev = scales::dollar(sum(net_price)))
```

```
##          tot_rev
## 1 $31,264,662
```


2. Consulting: full data

- What if we increased the price for those who submit scores?

```
# Although total admits declined
hypo %>%
  mutate(preds = predict(mLGFinal,newdata = hypo,type = 'response'))
%>%
  mutate(pred_class = ifelse(preds > .5,1,0)) %>%
  filter(pred_class == 1) %>%
  count()
```

```
##           n
## 1 1431
```

2. Consulting: full data

- Can we increase SAT scores to 1300?

```
ad %>%  
  filter(yield == 1) %>%  
  summarise(satAvg = round(mean(sat)))
```

```
##    satAvg  
## 1    1226
```

- Reduce price for those above 1300

2. Consulting: full data

- Can we increase SAT scores to 1300?
- Reduce price for those above 1300

```
hypo <- ad %>%
  mutate(net_price = ifelse(sat >= 1300,
                           net_price - 5000, net_price))

hypo %>%
  mutate(preds = predict(mLGFfinal, newdata = hypo, type = 'response'))
%>%
  mutate(pred_class = ifelse(preds > .5, 1, 0)) %>%
  filter(pred_class == 1) %>%
  summarise(satAvg = round(mean(sat)),
            tot_rev = scales::dollar(sum(net_price)),
            totAttend = n())
```

```
##   satAvg   tot_rev totAttend
## 1   1234 $28,294,452     1450
```

2. Consulting: full data

- Can we increase SAT scores to 1300?
- BUT need to make up losses

```
hypo <- ad %>%
  mutate(net_price = ifelse(sat >= 1300, net_price - 5000,
                           ifelse(sat < 1300, net_price +
2500, net_price)))

hypo %>%
  mutate(preds = predict(mLGFinal, newdata = hypo, type = 'response'))
%>%
  mutate(pred_class = ifelse(preds > .5, 1, 0)) %>%
  filter(pred_class == 1) %>%
  summarise(satAvg = round(mean(sat)),
            tot_rev = scales::dollar(sum(net_price)),
            totAttend = n())
```

```
##      satAvg      tot_rev totAttend
## 1      1236 $30,084,732      1414
```

2. Consulting: full data

- Looking across many different values? **Loops!**


```
toplot <- NULL
for(aid in c(2500,5000,7500,10000)) {
  for(charge in seq(2500,10000,by = 2500)) {
    hypo <- ad %>%
      mutate(net_price = ifelse(sat >= 1300,net_price - aid,
                                ifelse(sat < 1300,net_price +
charge,net_price)))

    tmp <- hypo %>%
      mutate(preds = predict(mLGFinal,newdata = hypo,type =
'response')) %>%
      mutate(pred_class = ifelse(preds > .5,1,0)) %>%
      filter(pred_class == 1) %>%
      summarise(satAvg = mean(sat),
                tot_rev = sum(net_price),
                totAttend = n()) %>%
      ungroup() %>%
      mutate(aid = aid,
             charge = charge)
```

```
toplot <- toplot %>%
```


2. Consulting: full data

```
toplot %>%  
  ggplot(aes(x = factor(aid), y = factor(charge), fill = satAvg)) +  
  geom_tile()
```




2. Consulting: full data

```
toplot %>%  
  ggplot(aes(x = factor(aid), y = factor(charge), fill = tot_rev)) +  
  geom_tile()
```



2. Consulting: full data

```
toplot %>%  
  ggplot(aes(x = factor(aid), y = factor(charge), fill = totAttend)) +  
  geom_tile()
```



Conclusion

- Remember the workflow:
 1. Train: `glm()`
 2. Predict: `predict()`
 3. Evaluate: `roc_auc()`
 4. Adjust: `net_price = ifelse()`
- Go to Brightspace and take the **15th** quiz
 - The password to take the quiz is ####
- **Homework:**
 - Problem Set 7 (due 2023-03-31 by 11:59PM)