

Problem Set 7

Classification Part 2

[YOUR NAME]

Due Date: 2023-11-10

Getting Set Up

Open RStudio and create a new RMarkdown file (.Rmd) by going to File -> New File -> R Markdown... . Accept defaults and save this file as [LAST NAME]_ps7.Rmd to your code folder.

Copy and paste the contents of this file into your [LAST NAME]_ps7.Rmd file. Then change the author: [YOUR NAME] (line 4) to your name.

All of the following questions should be answered in this .Rmd file. There are code chunks with incomplete code that need to be filled in.

This problem set is worth 10 total points, plus three extra credit points. The point values for each question are indicated in brackets below. To receive full credit, you must both have the correct code **and include a comment describing what each line does**. In addition, some questions ask you to provide a written response in addition to the code. Unlike the first two problem sets, some of the code chunks are totally empty, requiring you to try writing the code from scratch. Make sure to comment each line, explaining what it is doing!

You are free to rely on whatever resources you need to complete this problem set, including lecture notes, lecture presentations, Google, your classmates...you name it. However, the final submission must be complete by you. There are no group assignments. To submit, compile the completed problem set and upload the PDF file to Brightspace by midnight on 2023/11/10 Also note that the TAs and professors will not respond to Campuswire posts after 5PM on Friday, so don't wait until the last minute to get started!

Good luck!

ChatGPT Link [Optional]

*Copy the link to ChatGPT you used here: _____.

Question 0

Require tidyverse and tidymodels (for calculating AUC), and load the admit_data.rds (https://github.com/jbisbee1/DS1000_F2023/blob/main/Lectures/7_Classification/data/admit_data.rds?raw=true') data to an object called ad . (Tip: use the read_rds() function with the link to the raw data.)

```
require(tidyverse)
```

```
## Loading required package: tidyverse
```

```
## — Attaching packages — tidyverse 1.3.2 —
## ✓ ggplot2 3.3.6      ✓ purrr 0.3.4
## ✓ tibble 3.1.8       ✓ dplyr 1.0.10
## ✓ tidyr 1.2.1        ✓ stringr 1.4.1
## ✓ readr 2.1.2        ✓ forcats 0.5.2
## — Conflicts — tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag() masks stats::lag()
```

```
require(tidymodels)
```

```
## Loading required package: tidymodels
## — Attaching packages — tidymodels 1.0.0 —
## ✓ broom 1.0.1      ✓ rsample 1.1.0
## ✓ dials 1.0.0      ✓ tune 1.0.0
## ✓ infer 1.0.3      ✓ workflows 1.1.0
## ✓ modeldata 1.0.1  ✓ workflowsets 1.0.0
## ✓ parsnip 1.0.2    ✓ yardstick 1.1.0
## ✓ recipes 1.0.1
## — Conflicts — tidymodels_conflicts() —
## ✗ scales::discard() masks purrr::discard()
## ✗ dplyr::filter() masks stats::filter()
## ✗ recipes::fixed() masks stringr::fixed()
## ✗ dplyr::lag() masks stats::lag()
## ✗ yardstick::spec() masks readr::spec()
## ✗ recipes::step() masks stats::step()
## • Use tidymodels_prefer() to resolve common conflicts.
```

```
ad <- read_rds('https://github.com/jbisbee1/DS1000_F2023/blob/main/Lectures/7_Classification/data/admit_data.rds?raw=true')
```

Question 1 [3 points]

a. Compare a linear regression (`mLM <- lm(...)`) to a logit regression (`mLG <- glm(...)`) where you predict attendance (`yield`) as a function of the following X predictors:

- distance
- income
- sat
- gpa
- visit
- registered
- legacy
- net_price

Evaluate the model performance using `roc_auc` based on cross validation with 100 iterations, using an 80-20% split of the data [2 points].

b. Does the linear regression model or the logit perform better? [1 point]

```

set.seed(123)
# a.
cvRes <- NULL
for(i in 1:100) {
  inds <- sample(1:nrow(ad),size = round(nrow(ad)*.8),replace = F)
  train <- ad %>% slice(inds)
  test <- ad %>% slice(-inds)

  # Linear
  mLM <- lm(yield ~ distance + income + sat + gpa + visit + registered + legacy + net_price,train)

  # Logit
  mLG <- glm(yield ~ distance + income + sat + gpa + visit + registered + legacy + net_price,family = binomial(link = 'logit'),data = train)

  toEval <- test %>%
    mutate(predLM = predict(mLM,newdata = test),
           predLG = predict(mLG,newdata = test,type = 'response')) %>%
    mutate(truth = factor(yield,levels = c('1','0')))

  tmpLM <- roc_auc(toEval,truth = 'truth',estimate = 'predLM') %>%
    mutate(cvInd = i,
           algo = 'LM')

  tmpLG <- roc_auc(toEval,truth = 'truth',estimate = 'predLG') %>%
    mutate(cvInd = i,
           algo = 'Logit')

  cvRes <- cvRes %>%
    bind_rows(tmpLM) %>%
    bind_rows(tmpLG)
}

# b.
cvRes %>%
  group_by(algo) %>%
  summarise(auc = mean(.estimate))

```

```

## # A tibble: 2 × 2
##   algo    auc
##   <chr> <dbl>
## 1 LM    0.872
## 2 Logit 0.909

```

- Based on this analysis, the logistic regression model performs better, with an AUC of 0.91 versus 0.87 for the linear model.

Question 2 [3 points]

- Based on the result to question 1, choose the best classification algorithm and train it on the full data. [1 point]
- Calculate the specificity and sensitivity across different thresholds ranging from zero to one, and plot these as different colored lines. [1 point]
- What is the optimal threshold to balance the trade-off between sensitivity and specificity based on this plot? **HINT:** Use `geom_vline()` and test different `xintercept` values until you nail the intersection between the two lines. [1 point]

```
# a.
mFinal <- glm(yield ~ distance + income + sat + gpa + visit + registered + legacy + net_price, a
d,
              family = binomial(link = 'logit'))


# b.
ad <- ad %>%
  mutate(preds = predict(mFinal, type = 'response'))

topplot <- NULL
for(thresh in seq(0,1,by = 0.025)) {
  topplot <- ad %>%
    mutate(pred_attend = ifelse(preds > thresh,1,0)) %>%
    group_by(yield) %>%
    mutate(total_attend = n()) %>%
    group_by(yield,pred_attend,total_attend) %>%
    summarise(nStudents=n(),.groups = 'drop') %>%
    mutate(prop = nStudents / total_attend) %>%
    ungroup() %>%
    mutate(accuracy = sum((yield == pred_attend)*nStudents) / sum(nStudents)) %>%
    mutate(threshold = thresh) %>%
    bind_rows(topplot)
}

topplot %>%
  mutate(metric = ifelse(yield == 1 & pred_attend == 1,'Sensitivity',
                        ifelse(yield == 0 & pred_attend == 0,'Specificity',NA))) %>%
  drop_na(metric) %>%
  ggplot(aes(x = threshold,y = prop,color = metric)) +
  geom_line() +
  scale_x_continuous(breaks = seq(0,1,by = .05)) +
  geom_vline(xintercept = .615) + # c.
  labs(title = 'Sensitivity and Specificity by Threshold',
       subtitle = 'Model: Linear Regression',
       x = 'Threshold',
       y = 'Proportion Correct',
       color = 'Metric')
```

Sensitivity and Specificity by Threshold

Model: Linear Regression



- The optimal threshold is 0.615. This threshold value maximizes both sensitivity and specificity of the model.

Question 3 [4 points]

- How many students with SAT scores higher than 1300 are currently enrolled (`yield`)? How many students with SAT scores higher than 1300 are predicted to enroll according to our model? [1 point]
- What is the average SAT score and total tuition among enrolled students? [1 point]
- Reduce the net price (`net_price`) for students with SAT scores higher than 1300 by \$5,000. How many are now estimated to enroll? [1 point]
- What is the average SAT score among students predicted to enroll after adjusting the `net_price` ? What is the total tuition? [1 point]

```
# a. 314 high SAT students currently enrolled.
ad %>%
  count(yield,sat > 1300)
```

```
##   yield sat > 1300    n
## 1     0     FALSE 655
## 2     0      TRUE  29
## 3     1     FALSE 1152
## 4     1      TRUE  314
```

```
# 328 high SAT students predicted to enroll.
ad %>%
  mutate(preds = predict(mFinal,type = 'response')) %>%
  mutate(pred_attend = ifelse(preds > .615,1,0)) %>%
  count(pred_attend,sat > 1300)
```

```
##   pred_attend sat > 1300    n
## 1           0     FALSE 799
## 2           0      TRUE  15
## 3           1     FALSE 1008
## 4           1      TRUE  328
```

```
# b.
ad %>%
  filter(yield == 1) %>%
  summarise(avgSAT = mean(sat,na.rm=T),
            tuition = sum(net_price,na.rm=T))
```

```
##   avgSAT tuition
## 1 1225.941 30674149
```

```
# c.
hypo <- ad %>%
  mutate(net_price = ifelse(sat > 1300,net_price - 5000,net_price))

hypo %>%
  mutate(preds = predict(mFinal,newdata = hypo,type = 'response')) %>%
  mutate(pred_attend = ifelse(preds > .615,1,0)) %>%
  count(pred_attend,sat > 1300)
```

```
##   pred_attend sat > 1300    n
## 1           0     FALSE 799
## 2           0      TRUE  11
## 3           1     FALSE 1008
## 4           1      TRUE  332
```

```
# d.
hypo %>%
  mutate(preds = predict(mFinal,newdata = hypo,type = 'response')) %>%
  mutate(pred_attend = ifelse(preds > .615,1,0)) %>%
  filter(pred_attend == 1) %>%
  summarise(avgSAT = mean(sat,na.rm=T),
            tuition = sum(net_price,na.rm=T))
```

```
##      avgSAT  tuition
## 1 1238.878 25435144
```

- There are 314 students currently enrolled with SAT scores higher than 1300. Our model predicts that there should be 328 students enrolled who have SAT scores higher than 1300.
- The average SAT score is roughly 1226, and the total tuition is \$30.7m.
- After adjusting the price, the model predicts that 332 students with SAT scores greater than 1300 will enroll.
- After adjusting the price, the model predicts that the average SAT score will be 1239 and the total tuition will be \$25.4m.

Extra Credit [3 points]

- How high can you increase the average SAT score while maintaining current revenues, using only the `net_price` to induce changes? [1 point]
- Answer this question using a loop. [1 point]
- How does your answer change if you restrict the final `net_price` value per observation to be no lower than zero, and no higher than \$45,000? [1 point]

```
# a.
hypo <- ad %>%
  mutate(net_price = ifelse(sat >= 1300,0,45000))

hypo %>%
  mutate(preds = predict(mFinal,newdata = hypo,type = 'response')) %>%
  mutate(pred_attend = ifelse(preds > .615,1,0)) %>%
  filter(pred_attend == 1) %>%
  summarise(mean_sat = mean(sat),
            revenues = sum(net_price),
            nStudents = n())
```

```
##      mean_sat revenues nStudents
## 1 1238.299 35505000      1117
```

```

# b.
toplot <- NULL
for(i in c(2500,5000,7500,10000,15000,20000,25000)) {
  for(j in c(2500,5000,7500,10000,15000,20000,25000)) {
    for(realistic in c(TRUE,FALSE)) { # c.
      hypo <- ad %>%
        mutate(net_price = ifelse(sat >= 1300,net_price - i,net_price + j))

      # c.
      if(realistic) {
        hypo <- hypo %>%
          mutate(net_price = ifelse(net_price < 0,0,
            ifelse(net_price > 45000,45000,net_price)))
      }

      toplot <- hypo %>%
        mutate(preds = predict(mFinal,newdata = hypo,type = 'response')) %>%
        mutate(pred_attend = ifelse(preds > .615,1,0)) %>%
        filter(pred_attend == 1) %>%
        summarise(mean_sat = mean(sat),
          revenues = sum(net_price),
          nStudents = n()) %>%
        ungroup() %>%
        mutate(aid = i,
          charge = j,
          realistic = realistic) %>%
        bind_rows(toplot)

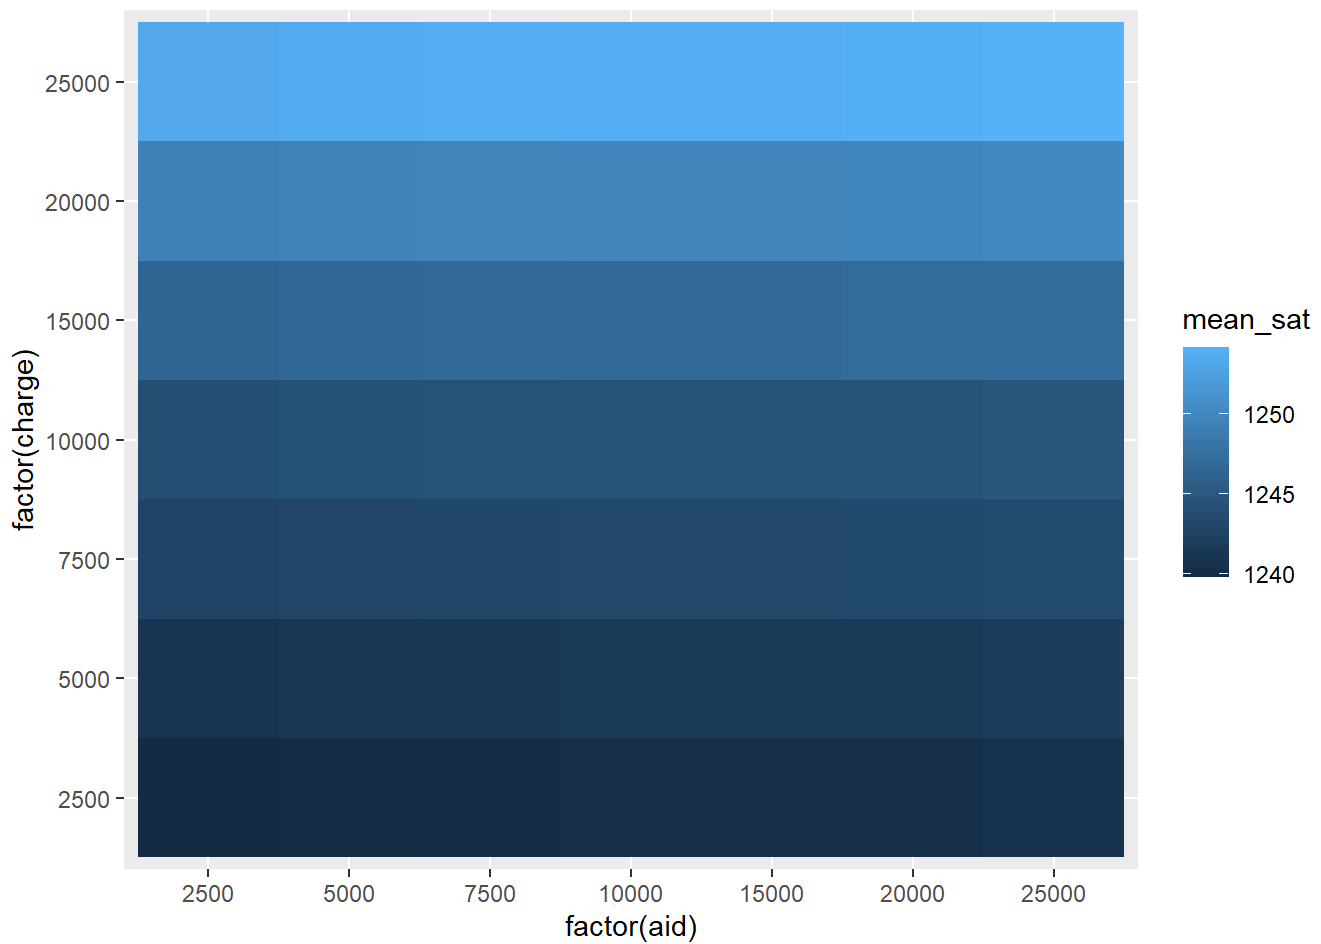
    }
  }
}

toplot %>%
  filter(charge == 25000)

```

##	mean_sat	revenues	nStudents	aid	charge	realistic
## 1	1254.217	29564701	1054	25000	25000	FALSE
## 2	1238.955	34619300	1172	25000	25000	TRUE
## 3	1254.021	31304701	1052	20000	25000	FALSE
## 4	1238.955	34619300	1172	20000	25000	TRUE
## 5	1253.841	33039701	1049	15000	25000	FALSE
## 6	1238.955	34619300	1172	15000	25000	TRUE
## 7	1253.841	34714701	1049	10000	25000	FALSE
## 8	1238.955	34825401	1172	10000	25000	TRUE
## 9	1253.792	35559701	1048	7500	25000	FALSE
## 10	1238.955	35387140	1172	7500	25000	TRUE
## 11	1253.605	36404701	1046	5000	25000	FALSE
## 12	1238.955	36025537	1172	5000	25000	TRUE
## 13	1253.358	37242201	1043	2500	25000	FALSE
## 14	1238.955	36696639	1172	2500	25000	TRUE



```
# b.  
toplot %>%  
  filter(!realistic) %>%  
  ggplot(aes(x = factor(aid), y = factor(charge), fill = mean_sat)) +  
  geom_tile()
```



```
toplot %>%  
  filter(!realistic) %>%  
  filter(revenues >= 30674149) %>%  
  arrange(desc(mean_sat))
```

##	mean_sat	revenues	nStudents	aid	charge	realistic
## 1	1254.021	31304701	1052	20000	25000	FALSE
## 2	1253.841	33039701	1049	15000	25000	FALSE
## 3	1253.841	34714701	1049	10000	25000	FALSE
## 4	1253.792	35559701	1048	7500	25000	FALSE
## 5	1253.605	36404701	1046	5000	25000	FALSE
## 6	1253.358	37242201	1043	2500	25000	FALSE
## 7	1249.891	32292349	1116	15000	20000	FALSE
## 8	1249.891	33967349	1116	10000	20000	FALSE
## 9	1249.842	34812349	1115	7500	20000	FALSE
## 10	1249.659	35657349	1113	5000	20000	FALSE
## 11	1249.416	36494849	1110	2500	20000	FALSE
## 12	1246.920	32184522	1175	10000	15000	FALSE
## 13	1246.871	33029522	1174	7500	15000	FALSE
## 14	1246.692	33874522	1172	5000	15000	FALSE
## 15	1246.454	34712022	1169	2500	15000	FALSE
## 16	1244.211	31416657	1228	5000	10000	FALSE
## 17	1243.977	32254157	1225	2500	10000	FALSE
## 18	1242.796	30985783	1253	2500	7500	FALSE

```
# c.  
toplot %>%  
  filter(realistic) %>%  
  ggplot(aes(x = factor(aid), y = factor(charge), fill = mean_sat)) +  
  geom_tile()
```



```
toplot %>%  
  filter(realistic) %>%  
  filter(revenues >= 30674149) %>%  
  arrange(desc(mean_sat))
```

##	mean_sat	revenues	nStudents	aid	charge	realistic
## 1	1239.146	31579121	1246	25000	15000	TRUE
## 2	1239.146	31579121	1246	20000	15000	TRUE
## 3	1239.146	31579121	1246	15000	15000	TRUE
## 4	1239.146	31785223	1246	10000	15000	TRUE
## 5	1239.146	32346961	1246	7500	15000	TRUE
## 6	1239.146	32985359	1246	5000	15000	TRUE
## 7	1239.146	33656460	1246	2500	15000	TRUE
## 8	1239.121	30834685	1279	5000	10000	TRUE
## 9	1239.121	31505787	1279	2500	10000	TRUE
## 10	1239.061	33356948	1211	25000	20000	TRUE
## 11	1239.061	33356948	1211	20000	20000	TRUE
## 12	1239.061	33356948	1211	15000	20000	TRUE
## 13	1239.061	33563050	1211	10000	20000	TRUE
## 14	1239.061	34124788	1211	7500	20000	TRUE
## 15	1239.061	34763185	1211	5000	20000	TRUE
## 16	1239.061	35434287	1211	2500	20000	TRUE
## 17	1238.955	34619300	1172	25000	25000	TRUE
## 18	1238.955	34619300	1172	20000	25000	TRUE
## 19	1238.955	34619300	1172	15000	25000	TRUE
## 20	1238.955	34825401	1172	10000	25000	TRUE
## 21	1238.955	35387140	1172	7500	25000	TRUE
## 22	1238.955	36025537	1172	5000	25000	TRUE
## 23	1238.955	36696639	1172	2500	25000	TRUE

- I can achieve an average SAT score of 1254 while maintaining revenues above \$3.07m if I reduce `net_price` by \$20,000 for students with SAT scores greater than or equal to 1300, and charging those with SAT scores lower than 1300 an additional \$25,000.
- If I restrict `net_price` to be realistic values (i.e., never going lower than zero and never greater than the maximum tuition of \$45,000), I can only improve average SAT scores to 1239 which is achieved by charging students with scores less than 1300 an additional \$15,000 and incentivizing those with scores 1300 or greater with any amount of money.